

# Introduction to Neural Networks


From Perceptron to Modern AI

Leila Kalhor  
Shahid Beheshti University

**ICTP, Physics Without Frontiers**  
October 2025

# Course Goals & Roadmap

- Understand the concept of an **artificial neuron**
- Grasp the **architecture of neural networks**
- See how ANNs **approximate complex functions**
- Survey key **applications** in AI




## Takeaway

By the end of this intro you should be comfortable with the building blocks and why they work.



# Artificial Neuron (Perceptron)

- Inputs  $\{x_i\}$  with **weights**  $\{w_i\}$  and **bias**  $b$
- Linear combination:  
$$z = \sum_i w_i x_i + b$$
- Output:  $y = f(z)$  where  $f$  is an **activation**
- Introduced by **Frank Rosenblatt** (1958)




## Key idea

Nonlinear activations let simple units build complex decision boundaries.



# A Short History (Visual Timeline)




## Context

Progress accelerated with data, compute (GPUs), and algorithms—enabling today's practical AI systems.



# Modern Applications (Visual Collage)

- **Computer Vision:** classification, detection, segmentation
- **NLP:** translation, question answering, sentiment analysis
- **Robotics/Control:** navigation, manipulation, autonomous driving




## Message

Neural networks are the *function approximators* behind many intelligent systems.



# Perceptron: Neuron Equation

- Linear part:  $z = \sum_i w_i x_i + b$
- Nonlinearity:  $y = f(z)$  (activation)
- Before  $f$ : **linear**; after  $f$ : **nonlinear**



## Key Idea

Simple units + nonlinear activations  $\Rightarrow$  expressive decision boundaries.




# Activation & Linear Separability

- If  $f$  is **linear**  $\Rightarrow$  the whole network remains linear.
- We need **nonlinear** activations (ReLU, Tanh, Sigmoid).
- Linear decision boundary works for linearly separable data.

## Takeaway

Nonlinearity unlocks complex decision surfaces.




# Perceptron Learning: Intuition

- Goal: **reduce loss** by nudging weights downhill.
- Update (intuition):  $w \leftarrow w - \eta \partial L / \partial w$
- Learning rate  $\eta$ : too big  $\Rightarrow$  oscillation; too small  $\Rightarrow$  slow.

## Mental Model

A ball rolling down the loss surface toward a minimum.






# Limitation: XOR is Not Linearly Separable

- Single perceptron fails on XOR — can't separate with one line.
- Motivation for **multi-layer** networks (MLPs).

## Message

Hidden layers combine simple boundaries to solve non-linear problems.




# MLP Architecture (Feedforward)

- Layers: Input  $\rightarrow$  Hidden (nonlinear)  $\rightarrow$  Output
- Nonlinear activations between linear layers
- Depth/width increase representational power

## Idea

Composition of simple units builds complex features.




# Forward Pass & Loss

- Forward:  $a^{(l)} = f(W^{(l)}a^{(l-1)} + b^{(l)})$
- Loss: Cross-Entropy (classification), MSE (regression)
- Goal: minimize loss over data

## Note

Normalization and stable outputs (e.g., Softmax) help training.




# Backpropagation: Intuition

- Compute output error  $\Rightarrow$  propagate **backwards**
- Chain rule links each weight to loss
- Update:  $W^{(l)} \leftarrow W^{(l)} - \eta \nabla_{W^{(l)}} L$


## Picture

Blue arrows forward (activations), **red dashed** arrows backward (errors).



# Gradient Flow: Vanishing vs Exploding

- Deep chains can **shrink** (vanish) or **blow up** (explode) gradients
- Remedies: ReLU-family activations, good initialization, normalization



## Visual

Arrow thickness indicates gradient magnitude across layers.




# Optimization & Learning Rate

- SGD / Mini-batch GD for efficiency
- Learning rate  $\eta$ : small = slow, large = unstable
- Practical: schedulers, momentum/Adam

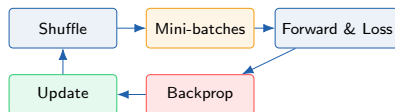
## Heuristic

Start with a conservative  $\eta$ , increase if stable; otherwise decrease.



# Training Loop & Mini-batching

- Shuffle data → split into mini-batches
- For each batch: forward → loss → backprop → update
- Monitor train/val metrics; early stopping on plateau




## Recipe

BatchNorm/LayerNorm, proper init, and regularization improve stability.



# Universal Approximation Theorem: Statement

- A feedforward network with one hidden layer and a **nonlinear** activation can approximate any **continuous** function on a compact domain, given enough neurons.
- UAT speaks about **expressive power**, not training ease.
- Depth is not required by the theorem, but often improves efficiency.



## Essence


With sufficient width, shallow networks are universal approximators.





# UAT Intuition: Building Functions from Simple Pieces

- Partition the input domain and stitch **simple pieces** together: steps or piecewise-linear segments (e.g., sums of ReLUs).
- Increasing the number of hidden units  $\Rightarrow$  finer partition  $\Rightarrow$  better approximation.
- Depth can reduce the number of units needed for a given accuracy.




## Picture

Target curve (blue) vs. shallow piecewise approximation (green).



# UAT: Misconceptions & Practical Limits

- **Expressivity  $\neq$  Learnability:** UAT does not guarantee training success.
- Data, optimization, and regularization control **generalization**.
- Shallow universality may require **many** neurons; depth can be **parameter-efficient**.
- Choice of activation matters (ReLU/tanh vs. saturating sigmoids).




## Takeaway

UAT explains *why* ANNs can represent complex functions—not *how* to train them well.



# Activation Zoo & Roles

- **Why:** Nonlinearity enables complex decision boundaries.
- **Sigmoid**  $(0, 1)$ : saturates; good for probabilities (binary).
- **Tanh**  $(-1, 1)$ : zero-centered; still saturates.
- **ReLU**  $\max(0, x)$ : simple, sparse, robust gradients.



## Takeaway

Hidden layers: ReLU-family as a strong default; outputs depend on task.




# Softmax & Cross-Entropy (Multi-class)

- **Softmax:**  $p_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$  converts logits to a probability simplex.
- **Cross-Entropy:**  $-\sum_i y_i \log p_i$  aligns predicted distribution with labels.
- Stable training: combine Softmax + CE; use log-sum-exp tricks in practice.

Logits

Probabilities



## Message

Use Softmax at the output for single-label multi-class problems.






# Practical Tips & Pitfalls

- **Vanishing gradients:** Sigmoid/Tanh saturate  $\Rightarrow$  use ReLU-family or normalization.
- **Dead ReLU:** neurons stuck at  $x < 0$ ; mitigate with LeakyReLU/ELU/GELU.
- **Choices:**
  - Hidden: ReLU / LeakyReLU (safe defaults)
  - Output: Softmax (multi-class), Sigmoid (multi-label), Linear (regression)


## Rule of Thumb

Start with ReLU (or LeakyReLU), change only if gradients/accuracy suggest otherwise.



# What is Overfitting?

- Model fits **noise** or idiosyncrasies of training data
- Training loss  $\downarrow$  while **validation loss** eventually  $\uparrow$
- Poor **generalization** to unseen data



## Signal

Growing gap: train vs. validation metrics as epochs increase.




# Regularization Toolbox

- **Weight Decay (L2):** penalize large weights
- **Dropout:** randomly deactivate units during training
- **Early Stopping:** stop at best validation performance
- (Also: Data Augmentation, Norm layers, Smaller models)


## Goal

Reduce variance without adding too much bias.



# Validation & Early Stopping

- Split data: **Train / Validation / Test**
- Monitor validation loss/accuracy each epoch
- Stop when validation no longer improves (**patience**  $k$  epochs)



## Outcome

Prevents over-training past the generalization sweet spot.






# Symptoms & Sanity Checks

- **Large gap** train vs. val/test metrics
- Highly complex model vs. small dataset
- Unstable training, high variance across runs
- Fixes: regularization, more data/augmentation, simpler model


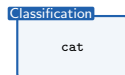
## Quick Checks

Shuffle properly, hold-out a test set, verify labels/leakage.



# Applications: Computer Vision

- **Classification:** image-level labels
- **Detection:** bounding boxes for objects
- **Segmentation:** pixel-level understanding
- Pipelines: data → augment → CNN/MLP head → metrics



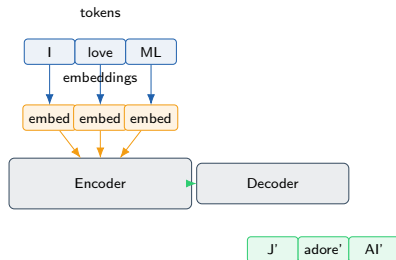
## Impact

From medical imaging to autonomous driving and retail analytics.



# Applications: Natural Language Processing

- **Machine Translation, Sentiment Analysis, Question Answering**
- Tokenization → Embeddings → Encoder/Decoder → Output
- Losses: Cross-Entropy, Label Smoothing; decoding: Greedy/Beam



## Note

Context modeling is key;  
attention/transformers extend MLP  
fundamentals.




# Applications: Tabular/Time-Series & Robotics/Control

- **Regression/Forecasting:** prices, demand, sensors
- **Anomaly Detection:** monitoring, security
- **Control:** policy/value approximation for decision making

## Pattern

Learned representations beat manual features when data is sufficient.



# Summary & Next Steps

- **Core:** Perceptron → MLP/Backprop → UAT
- **Practice:** Activations, optimization, regularization
- **Applications:** Vision, NLP, Tabular/Control
- **Next:** hands-on demo + try hyperparameter tweaks



- ReLU/Softmax choices
- LR, batch size, schedulers
- L2/Dropout/Early stop

## Key Message

Neural networks = powerful function approximators; training craft makes them useful.



# Thanks! *Backprop of Gratitude*

## Training Stats

- final loss:  $\approx 0$  (ish)
- smiles accuracy: 99.9% (val)
- optimizer: Adam (caffeinated)
- batch size: **you**
- regularizer: coffee & great questions
- epoch: until **Q&A** converges

**Thank You!**  
gradient of gratitude ↗

