Nayara Fonseca
University of Oxford

Big-science choices (such as the next collider) carry billion-dollar stakes and sharp disagreements. I will present a lightweight AI-debate prototype that makes these disagreements explicit and measurable. Two language-model "debaters" argue for alternative particle-physics roadmaps while a third, neutral model acts as a transcript-only judge. To reduce bias, each match uses the same source summaries, randomized order on every turn, alternating opener, balanced stance labels, and fixed prompts with token caps. We evaluate two measures: (i) stance strength (pooled win rates of the strategies) and (ii) debater strength (Elo-style ratings). Preliminary results show small, overlapping differences across stances and debaters. The goal is not to replace expert review, but to provide a transparent way to organize evidence and compare arguments.